

Introduction to Azure Storage

Big data storage technologies

- Azure Storage blobs
- Azure Data Lake Store
- Azure Cosmos DB
- HBase on HDInsight

Key Selection criteria

- size of dataset
- performance requirements (Horizontal or Vertical scaling)
- data types to be ingested
- where will data be consumed
- How will data be accessed,
- Data redundancy and availability

Big data Storage Technology

SOL Database

Purpose Transaction processing

Data Capture Single source

Benefits Consistent data that can handle complex queries

Data quality Organized and consistent

SOL Data Warehouse

Purpose Multiple source reporting and queries

Data Capture Multiple relational sources

Benefits Complex data that can handle complex queries through MPP

Data quality Organized and consistent

Data Lake Storage

Purpose Dynamic big data storage at any speed

Data capture Any number of sources

Benefits Different data types that can change at will

Data quality Completely dependent upon Rules of Engagement (ROE)

File Storage

Collection of objects and files organized into a hierarchy of directories and subdirectories

- file systems that can be used by anyone

Cosmos DB

Purpose Global distribution

Cosmos DB

Purpose **Global distribution**

Data capture Multiple geo-replicated sources

Benefits Multi-modal, elastic, global-replication

Data quality organized and consistent

Cosmos DB is non-relational database

Global distribution and speed \Rightarrow Cosmos DB

Consistency and integrity \Rightarrow SQL DB / data warehouse

Choosing the right data store

General considerations

- Functional requirements
- Non-functional requirements
- Management and costs
- Security
- DevOps

DP-201 scope

- RDBMS
- Key/Value stores
- Document database
- Column-family databases
- Graph databases

- SQL database
- SQL data warehouse
- Cosmos DB
 - * SQL API
 - * Gremlin API
 - * MongoDB API
 - * Cassandra API
 - * Table API

RDBMS

Two dimensional tables with rows and columns

structured schema on write

uses SQL

Incorporates ACID model

Use cases:

- strong consistency guarantees are important
- schema provides opportunity for complex queries

future applications:

Azure applications:

Azure SQL database

Azure SQL data warehouse

Key-value stores:

- Simple databases that uses an associative array.
↳ Key is associated with one value in collection
- Query, insert and delete operations

Typical use cases

- user profiles
- Product recommendations
- Session management

Azure applications:

- Cosmos DB
- Redis cache

Document Database

deux key-value store

documents can be encoded in XML, JSON, YAML etc.

schema provides opportunity for complex queries

free form storage system

Azure Applications

- Cosmos SQL API
- Cosmos MongoDB API

Column family databases:

- Similar to RDBMS. Each table has rows and columns.

Groups related columns into families

Duplicates are possible

Joins are out

Azure Applications:

- HBase in HDInsight
- Cosmos Cassandra API

Graph databases:

Nodes = Entities

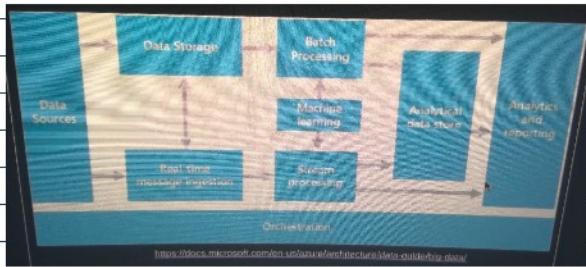
Edges = Relationships

Azure applications:

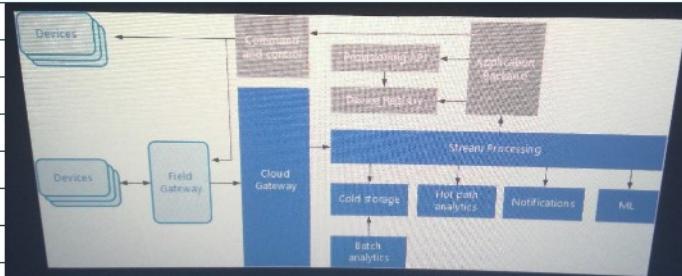
Cosmos Gremlin API

Big data Architecture





Machine Learning
Analytical data store
↳ HBase
↳ SQL Data Warehouse



Capability Matrix:

File storage capabilities

| Capability | Azure Data Lake Store | Azure Blob Storage containers |
|--------------------------------|--|---|
| Purpose | Optimized storage for big data analytics workloads | General purpose object store for a wide variety of storage scenario |
| Use cases | Batch, streaming analytics, and machine learning data such as log files, IoT data, click streams, large datasets | Any type of text or binary data, such as application back end, backup data, media, storage for streaming, and general purpose data |
| Structure | Hierarchical file system | Object store with flat namespace |
| Authentication | Based on Azure Active Directory Identities | Based on shared secrets Account Access Keys and Shared Access Signature Keys, and role-based access control (RBAC) For account level authorization use Account Access Keys. For account, container, or blob authorization use shared Access Signature Keys |
| Authentication protocol | OAuth 2.0 Calls must contain a valid JWT token issued by Azure Active Directory | |
| Auditing | Available | Available |
| Encryption at rest | Transparent, server side | Transparent, server side, client side encryption. |
| Developer SDKs | .NET, Java, Python, Node.js | .NET, Java, Python, Node.js, C#, Ruby |
| Analytics Workload Performance | Optimized performance for parallel analytics workloads, High Throughput and IOPS | Not optimized for analytics workloads |
| Size limits | No limits on account sizes, file sizes or number of files | Specific limits exists |

| | | |
|----------------|--|-----------------------------|
| Geo-redundancy | Locally redundant (LRS), globally redundant (GRS), read-access globally redundant (RA-GRS), zone-redundant (ZRS) | LRS GRS ZRS RA-GRS |
|----------------|--|-----------------------------|

Nosql database capabilities:

| Capability | Azure Cosmos DB | HBase on HDInsight |
|-------------------------------------|--|---|
| Primary database model | Document store, graph, key-value store, wide column store | Wide column store |
| Secondary Indexes | Yes | No |
| SQL language support | Yes | Yes (using Phoenix JDBC driver) |
| Consistency | Strong, bounded staleness, session, consistent prefix, eventual | Strong |
| Native Azure Functions integrations | Yes | No |
| Automatic global distribution | Yes | No, HBase cluster replication can be configured across regions with eventual consistency. |
| Pricing model | Elastically scalable request units (RU) charged per-second as needed, elastically scalable storage | Per-minute pricing for HDInsight cluster (horizontal scaling of nodes), storage |

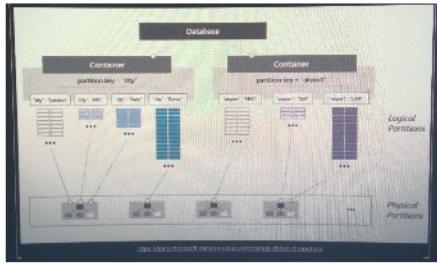
Data partitioning in Azure

Partitioning in Cosmos DB

- used to scale individual containers to achieve performance.
- Logical partitions are based on partition keys
- Physical partitions have components that can vary

Logical partitions

- automatically formed based on partition keys
- Item index is a combination of PartitionKey and item ID



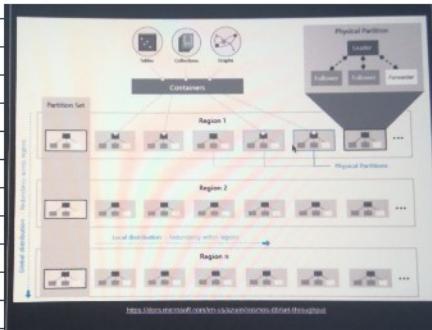
Logical partitions - building a Partition Key

Even distribution

- Partition key is property that will exist on every object
- Every document with the same PartitionKey belongs to the same logical partition
- PartitionKey decides the placement of documents.

Building a great PartitionKey

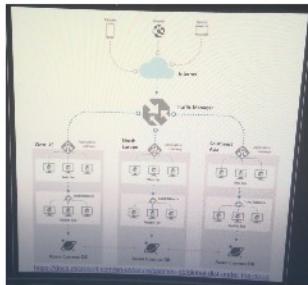
- Anticipate top queries
- Avoid fans
- Choose well balanced Keys
- Everything is subject to change
- Keys are immutable, so choose wisely
- 10 GB data maximum



Design for Global Distribution in Azure

Importance of global distribution in Cosmos DB

- **Unlimited elastic** write and read scalability
- **99.999%** read and write availability worldwide
- Guaranteed read and write **< 10 milliseconds** at the **99th percentile**



If a region goes down, Cosmos DB Traffic Manager automatically serves traffic from other region.

\$ Cosmos DB Pricing \$

Per container → Throughput Storage
 $T \times N$ $T \times (N+1)$ RU/s

Single region Multi region writes
 Write writes
 $T \rightarrow \text{Throughput}$
 $N \rightarrow \# \text{ regions}$

- \$ 25 / GB per month / per region
- No max storage limits

Consistency in Cosmos DB

① Strong Consistency

- Highest consistency
- lowest availability
- High cost

② Bounded Staleness

- lag between updates (e.g. scoreboard updated frequently without interrupting the game)

③ Session consistency

- Default for Cosmos DB
- Game A Game B games

Consistency across single game
 Lag for the other game

- same session consistency
- Cross session lag

④ Consistent prefix

Game A and Game B both are updated parallelly. Final scores are eventually consistent.

⑤ Eventual Consistency

- weakest consistency
- lowest resources

Cosmos DB API

Chosen at account creation

⚠ Cannot be modified later

① Core (SQL) API

- default
- SQL like
- JavaScript type, expression evaluation, and function invocation
- provides document data model

② MongoDB API

- supports MongoDB Wire protocol
- full compatibility with existing MongoDB app
- Built in find method

③ Cassandra API

- Columnar data model
- can define data schema upfront
- Cassandra Query Language (CQL)

④ Azure Table API

- Great for table storage with advanced requirements
- No index management required
- Key-value data model
- New projects, shift to SQL API

⑤ Gremlin API

- Columnar data model
- can define data schema upfront
- Cassandra Query Language (CQL)

■ New projects, shift to SOL API

⑤ Gremlin API

- Graph data model
- Edges & nodes or vertex

DISASTER RECOVERY & HIGH AVAILABILITY

High Availability

- multiple regions at least $N \times 4$ copies
- single region 99.99 SLA
- multi region
 - single writes 99.99 W, 99.999 R
 - multiple writes 99.999

Need to enable automatic failover
to avoid conflicts when the
region comes back online due to
unreplicated data.



Disaster Recovery

Online backup & data restore

- Auto backup every 4 hours → Only 2 latest snapshots are available for 30 days
- Blob storage
- Azure Cosmos DB for data

⚠️ No additional costs for automated backups

Additional Backups

- Auto geo-redundant storage
- Configure multi-region

- ADF or Cosmos DB changed feeds can be used for custom backup needs

Data Lake Gen 2

Best of Gen 1 +
blob storage

- Hot, cold and Archive storage (Access Tier)
- Redundant
- HDFS compatibility
- Hierarchical file system
- Elastic, scalable environment with enhanced performance

PARTITIONING TABLES

Table Partitions

- Smaller groups of data
- Improve maintenance
- Query performance

DATA DISTRIBUTION

HASH Distribution

- Row based on values
- High performance queries large tables

Replicated tables

- Full copy
- Small tables with performance

Round Robin

- Rows distributed evenly
- Random distribution
- Faster queries
- Improves load speed

⚠ Used for staging tables

⚠ Does not require data movement

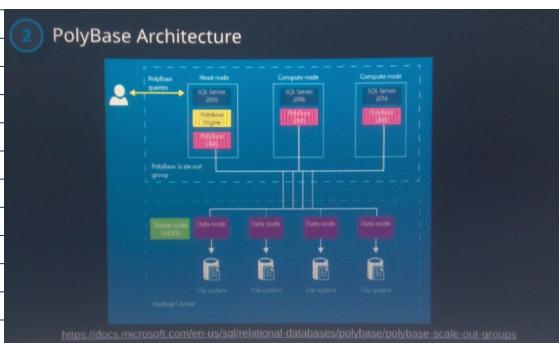
INGESTION WITH POLYBASE

Polybase

- Query across relational and Hadoop data
- Bypass traditional ETL for ELT
 - Perform transformations in DW
 - Distribute transformations across multiple databases

ELT steps

- Extract source data into text files
- Land data into Azure Blob or ADLS
- Prepare data for loading
- Load data into staging with PolyBase
- Transform
- Insert into production tables

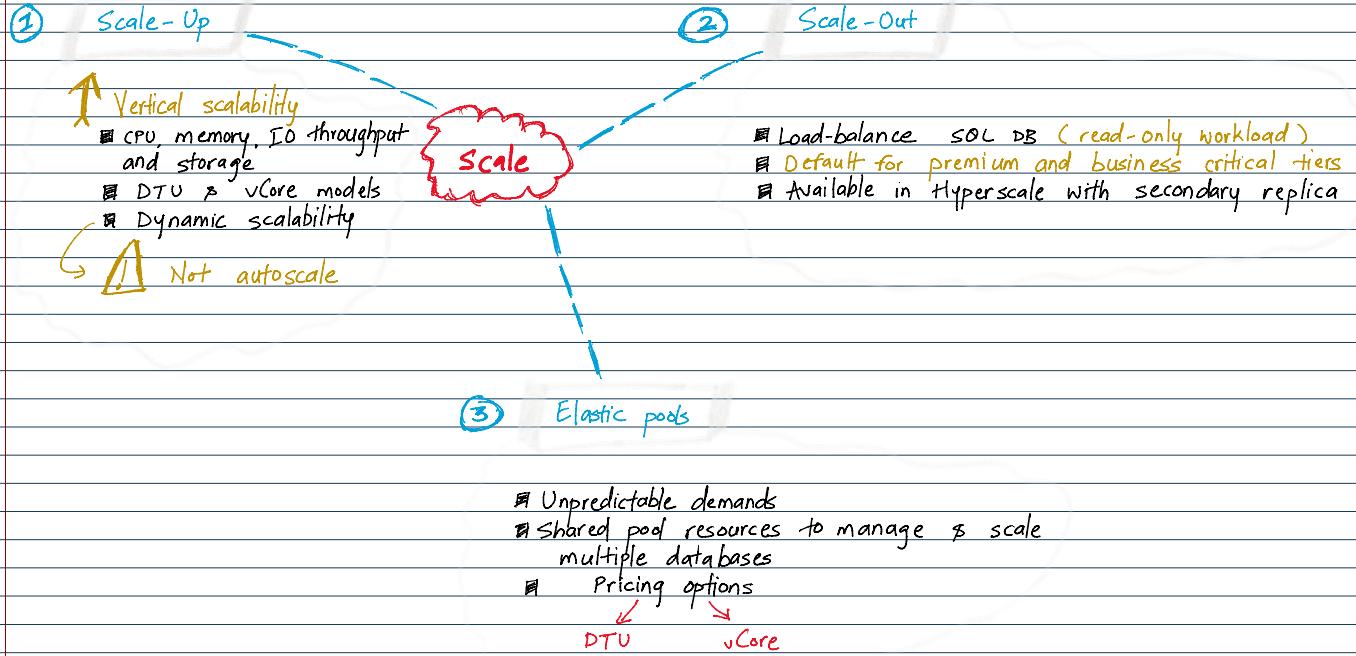


⚠ CTAS : optimized for bulk loads

```

CREATE TABLE [www].[dimension_City]
WITH
C
DISTRIBUTION = REPLICATE,
CLUSTERED COLUMNSTORE INDEX
)
AS
SELECT * FROM [ext].[dimension_City]
  
```

Design for Scale in SQL DB

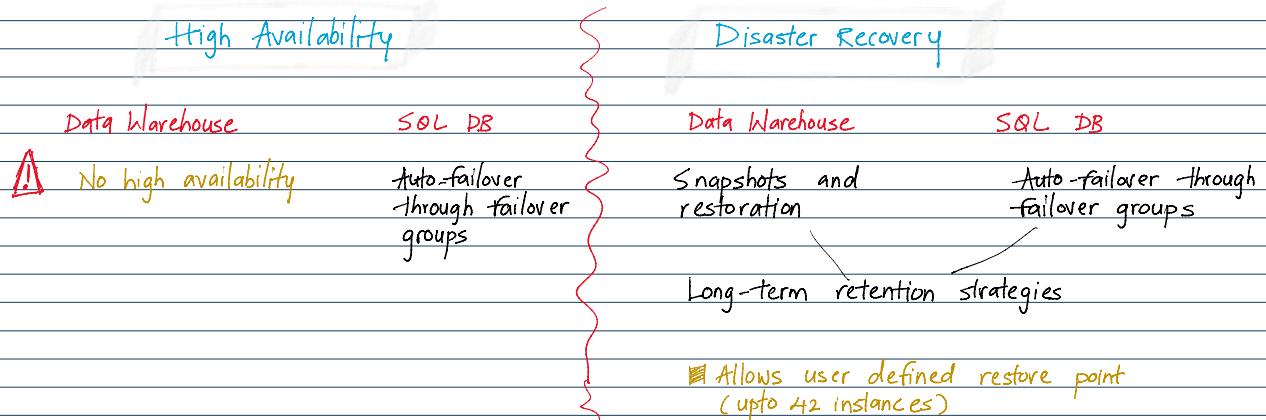


Scale SQL DW

DW scale or pause

- Scale up during heavy demand.
- Pause to cut cost.
- Modify with GUI, Powershell or TSOL.

Disaster Recovery & High Availability

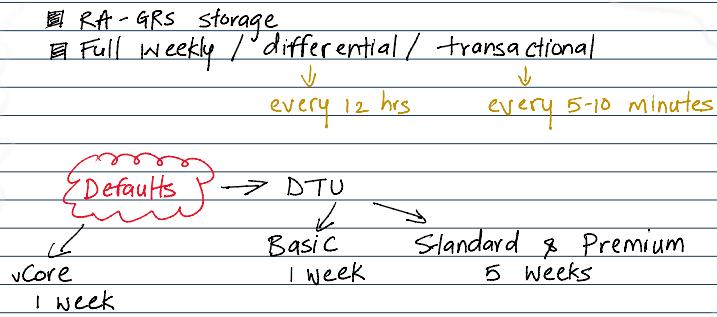


Allows user defined restore point
(upto 42 instances)

Default retention period → 7 days

SQl DB Automated Backup

Automated Backups



Long-Term Retention

Leverages PITR - blob for long-term storage

Eg.
■ W=0, M=0, Y=2, WeekOfYear=2

■ W=0, M=4, Y=0

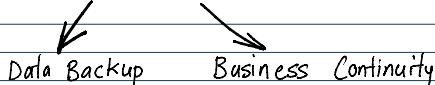
■ W=3, M=12, Y=7, WeekOfYear=10

Point-in-time Restore (PITR) max 14 days

Disaster Recovery Strategy

Protects organization from unacceptable negative events following a disaster.

2 core elements



Recovery Time Objective (RTO)

Maximum time before system is brought back online

Recovery Point Objective (RPO)

Acceptable gap of data lost when restoring backups

Data Ingestion

① Data Factory

- Orchestration Service
- ETL / ELT
- Core functions
 - Ingest
 - Data flow
 - Automate
 - Monitor
- Copy function

② Storage Explorer & Az Copy

- Storage explorer
 - Manual file or folder movement
- Az copy
 - Blob or files from storage account
 - local computers or other storage accounts

③ Apache Swoop & others

- Apache Swoop
 - Data movement tool into Hadoop HDFS
 - ETL tool

Data Processing

ELT

- | | | |
|---|---|---|
| ■ Self service | ■ Time | ■ Cost and complexity |
| <ul style="list-style-type: none"> • Data Lakes and zones • Raw data accessed by multiple thought leaders | <ul style="list-style-type: none"> • Accessibility • Maintenance • Load and transfer | <ul style="list-style-type: none"> • Data quality • Data management • Cost <ul style="list-style-type: none"> * Storage * Movement * Transformation * Resources |

Processing with Databricks

Terminology

- Cluster
- Workspaces
- Notebooks
- Cells
- Libraries

Key Features

- Interactive workspace for exploration and visualization
- Core API
 - R
 - SQL

- Notebooks
- Cells
- Libraries
- Tables

- R
- SQL
- Python
- Scala
- Java

- Streaming and batch processing
- Dynamic scaling and secure integration

Batch Processing using Data Factory

Core concepts

Pipeline

- Logical group of activities

Datasets

- Data structures within data stores

Activity

- Processing step in pipeline
- 3 types

Linked services

- Connections to data

↓
Data movement

↓
Data transformation

↓
Control



Inactive vs active

Pipeline orchestration and execution

Data flow execution and debugging

Data factory operations (e.g. creating pipelines & monitoring)

Stream Processing in Azure

Azure Stream processing use cases

- Never-ending streams of data
- Extreme precision is not needed
- IOT

Complex big data solutions involve stream and batch processing

Disaster Recovery

- Handle lost data
 - Windowing
 - Early and late arriving events

Input

Stream Analytics

Query

- Defines and connects inputs

Event hubs
IOT Hub
Blob storage

Output

- Build and test stream analytics queries
- Estimate pricing

Store and save data
Configure and validate receiving service

Data Lake, Cosmos DB, SQL DB / DW
Service Bus, Azure Functions, Event Hubs
Power BI



Stream Analytics jobs can be deployed on Cloud or Edge device

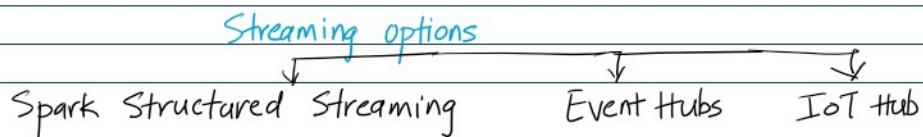


Pricing is based on streaming units

\$\$ Pricing is based on streaming units

Event ordering allows handling of late and out of order events

Streaming in Databricks



Optimize Stream Analytics

① Streaming Units

- Computing resources allocated to job
- Review SU% \Rightarrow 80% should be redline
- Start with 6 SUs for queries not using PARTITION BY

② Parallelization

- Partitioning key for optimization
 - Inputs are already partitioned
 - Outputs need to be partitioned
- Events should go to same partition of input
- Use PARTITION BY in all steps
- Input partition must equal output partition

Optimize Streaming in Databricks

① Recover query failures

② Spark scheduler pools

① Recover query failures

- Enable checkpointing
- Restart jobs on failures
- Recover after changes

② Spark scheduler pools

- Fair scheduler pools
 - Group jobs into pools by weight

③ Optimize configuration

- Use updated DBR
- Compute optimized instances
- shuffle partitions → 1-2x cluster number

④ Watermark policies

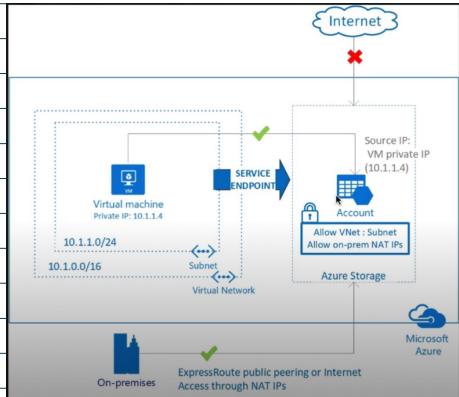
⑤ Auto Scaling

Streaming job config

- Max concurrent jobs → 1
- Timeout → None
- Retries → Unlimited

Secure Endpoints

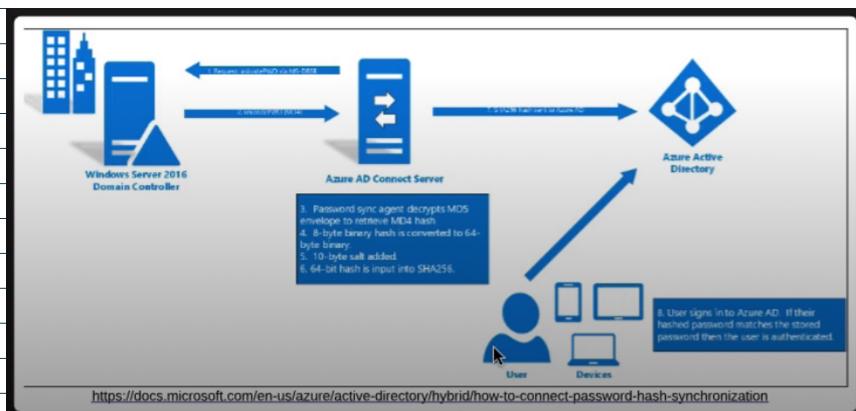
- Virtual Network (VNet) service endpoint
- Network extender to include Azure services



Private Link

- Creates private endpoint for Azure services
- Creates private IP restricting the traffic to Azure backbone

Authenticate with Azure Active Directory

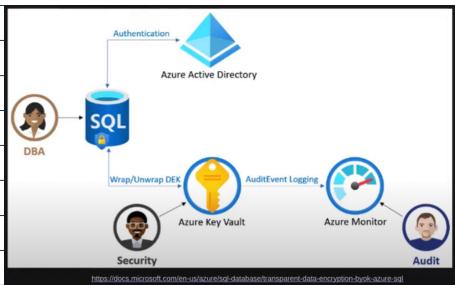


Authenticate with Access keys

Authenticate with Access keys

Azure Key Vault

- Securely store and access
 - Tokens
 - Passwords
 - Certificates
 - API keys
- Centralize storage of secrets
- Monitor access



Authenticate with Shared Access Signature

Shared Access Signature

- Secure delegated access

① User delegation SAS

⚠ Blob storage only

SAS types

② Service SAS

■ Storage Account Key
■ Azure Storage services only

③ Account SAS

■ Storage Account Key
■ Delegates to one or more storage services

Storage Account Key
Delegates to one or more storage services

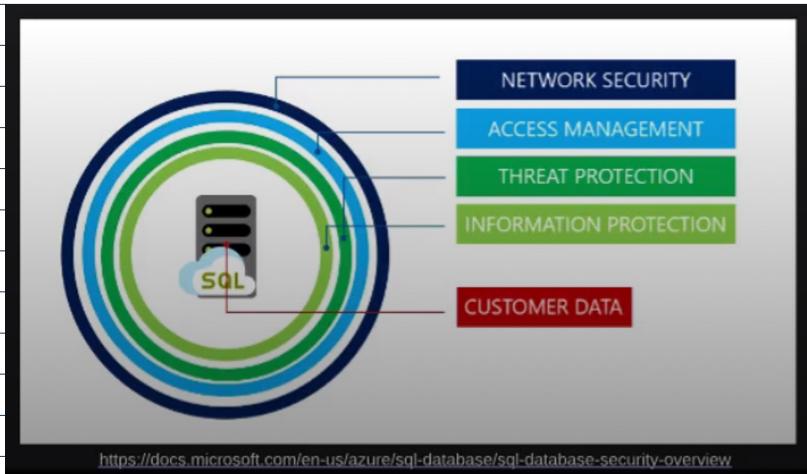
Defense in Depth

CIA

Confidentiality
Integrity
Availability

Security Layers

- Physical security - Building and computing hardware
- Identity and access - MFA
- Perimeter - Firewall
- Network - Network connectivity
- Compute - VM ↗ other compute resources
- Application
- Data - Dynamic data masking

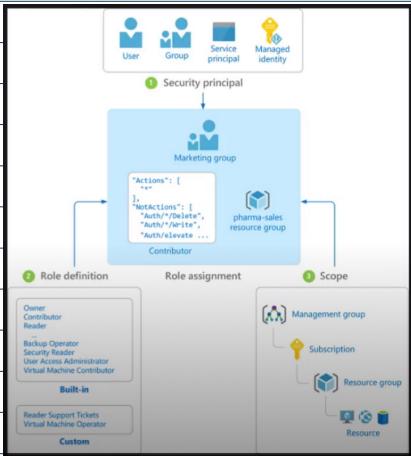


RBAC

■ Role Assignment

- Enforces permissions
- Types

↓ ↓ ↓
Security Principal Role definition Scope



■ Data Encryption in Transit & at Rest

Available for SaaS, PaaS, IaaS

Transparent Data Encryption

Available for

- SQL Server
- Azure SQL DB
- Azure Synapse

Uses

- AES and Triple Data Encryption standard
- Data encryption key

Cosmos DB Database Encryption

■ Encrypted by default

■ Uses secure key storage systems, encrypted networks and cryptographic APIs

encrypted networks and cryptographic APIs

Data Lake Encryption

- Encrypted by default
- MS managed keys, can be BYOK
- 3 keys
 - Master Encryption Key (MEK)
 - Data Encryption Key (DEK)
 - Block Encryption Key (BEK)

Encrypting Data in Motion

- TLS / SSL
 - Transport Layer Security → Between cloud and customer
 - Strong authentication
 - Message privacy
 - Integrity
 - Perfect Forward Secrecy (PFS)
 - ↳ Protects data between customer's client systems and cloud

- Shared Access Signatures
 - Delegate access to Azure Storage object

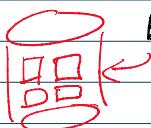
- Data Lake
 - Data in transit always encrypted
 - HTTPS is the only protocol available for REST interface

SQl Database Auditing

Use cases

- Retain → Audit trails
- Report → Event reporting
- Analyze → Spot trends and unusual activities

Audit policy



■ Server-level

- Applies to all existing and new databases
- Recommended approach



■ Database-level

- Specific database auditing
- Automatically enabled for secondary databases when geo-replication is active

Data Masking

Dynamic Data Masking

■ Real-time data masking

■ Limits sensitive data exposure to non-privileged users by masking

■ Types

- Default
- Credit card
- Email
- Random
- Custom

Use cases

- Protect non-production data
- Protect against insider threats
- Regulatory requirements

Archiving Strategy

Archiving Strategy

Azure Blob Storage

Hot

- Frequently used data
- Fast and expensive
 - * Low read and write cost
 - * High storage cost

Cool

- 30 days
- short-term backup and DR
 - * 1/2 cost compared to hot
 - * 2x cost to read and write

Archive

- 180 days
- Affordable[†] if we don't move data
- Long time to retrieve data
- Compliance and regulatory

Data Retention Policy

Tape vs Archive

- Convenience → cloud based vs physical
- Discovery → GDPR & other regulatory requirements
- Protection → Physical vs cloud
- Maintenance → Tape corruption, cloud always secure

Cool vs archive

Cost

- Storage
- Data movement
- Early deletion fees

Accessibility

Lifecycle Management

Lifecycle Management

- Rules for automating blob tier movement
- Policies to drive movement of data:
 - from hot to cool
 - from cool to archive
 - delete blobs
 - delete snapshots

Azure Data Security Best Practices

Key Vault

- Safeguard Keys, certificates and secrets
- Control access points

Protect data at rest

- Automated (ADLS, SQL DB, DW)
- TDE

Data in transit

- Endpoints
 - HTTPS
 - Azure storage
 - ExpressRoute

Active Directory

- Authorization
- Authentication
- RBAC
 - Security principal
 - Role
 - Definition